

# GenAI System Design

## Interview Guide


### Production Level RAG Architecture for Code Search

By Rajesh Srivastava

---

#### Interview Question

Design a scalable, low-latency code intelligence platform that can understand developer questions and fetch related code from company's 100K+ repositories



- Reference Code (Lightweight version): [Link](#)
- Latest LLMs, Embedding Models, Vector DBs comparison doc: [Link](#)

## Purpose of This Guide

This guide is structured to help you approach such system design interviews with clarity.

With minor tweaks, you can reuse the same framework for many RAG-based interview questions. Follow this approach:

**Phase 1: Clarify Requirements (2-3 minutes)** Start by asking clarifying questions to understand scope and constraints. Don't jump into the solution immediately.

**Phase 2: High-Level Design (5-7 minutes)** Draw the architecture overview, explain the two main pipelines (indexing and query), and identify major components.

**Phase 3: Deep Dive (15-20 minutes)** Based on interviewer interest, dive deep into specific areas. This document provides detailed sections on each component.

**Phase 4: Trade-offs & Optimizations (5-10 minutes)** Discuss alternative approaches, scaling strategies, and how you'd debug issues.

# Table of Contents

<b>Interview Opening Script .....</b>	5
<b>Key Clarifying Questions to Ask .....</b>	5
<b>Problem Context &amp; Why this is Hard .....</b>	6
<b>System Requirements (state these explicitly) .....</b>	6
<b>Architecture Overview .....</b>	7
<b>Section 1: Indexing Pipeline (Offline) .....</b>	8
<b>1.1 Code Parsing .....</b>	8
<b>1.2 Chunking Strategy .....</b>	8
<b>1.3 Embedding Generation .....</b>	9
<b>1.4 Storage Architecture.....</b>	10
<b>1.5 Incremental Indexing.....</b>	11
<b>Section 2: Query Pipeline (Online).....</b>	12
<b>2.1 Query Understanding.....</b>	12
<b>2.2 Multi-Stage Retrieval.....</b>	13
<b>Stage 1: Vector Similarity Search .....</b>	13
<b>Stage 2: Metadata Filtering &amp; Boosting .....</b>	13
<b>Stage 3: Cross-Encoder/flashrank Reranking.....</b>	13
<b>2.3 Response Generation .....</b>	14
<b>2.4 Caching Strategy .....</b>	14
<b>Section 3: Evaluation Metrics.....</b>	15
<b>3.1 Retrieval Metrics .....</b>	15
<b>3.2 Generation Metrics .....</b>	15
<b>3.3 Latency Metrics .....</b>	15
<b>Section 4: Scaling Strategy .....</b>	16
<b>4.1 Horizontal Scaling .....</b>	16
<b>4.2 Database Sharding.....</b>	16
<b>4.3 Cost Breakdown.....</b>	16
<b>Section 5: Debugging Poor Results .....</b>	17
<b>Step 1: Measure Current Performance .....</b>	17
<b>Step 2: Categorize Failure Modes.....</b>	17
<b>Step 3: Fix Biggest Issue .....</b>	17
<b>Step 4: Monitor Improvement .....</b>	17
<b>Section 6: Advanced Optimizations .....</b>	18
<b>6.1 Hybrid Search (Vector + BM25).....</b>	18
<b>6.2 Personalization .....</b>	18

<b>6.3 Continuous Learning .....</b>	18
<b>Interview Talking Points (At least Memorize these) .....</b>	18
<b>When discussing scale.....</b>	18
<b>When discussing accuracy.....</b>	18
<b>When discussing debugging.....</b>	18
<b>When discussing cost .....</b>	19
<b>Quick Reference Numbers (For reference) .....</b>	19
<b>Production Checklist.....</b>	20
<b>Infrastructure.....</b>	20
<b>Monitoring.....</b>	20
<b>Quality .....</b>	20
<b>Security.....</b>	20

# Interview Opening Script

When the interviewer presents this problem, start with:

---

Before I dive into the solution, let me make sure I understand the requirements correctly. We are building an internal code search tool (essentially a **Perplexity for Code**), where developers can ask natural language questions like 'Where is the payment validation logic?' and get back the exact function and file location. Let me clarify a few things...

---

## Key Clarifying Questions to Ask

1. **Scale:** How many repositories are we talking about? Are these all-internal repos or do we include open-source dependencies?
  2. **Users:** How many developers will use this? What is the expected QPS (Queries Per Seconds)?
  3. **Latency:** What is an acceptable response time? Sub-second, or can we tolerate a few seconds for complex queries?
  4. **Accuracy vs Speed:** Is it more important to always return the best result, or to return a good enough result quickly?
  5. **Languages:** Are we supporting all programming languages or a subset?
  6. **Freshness:** How quickly should new code changes appear in search results? Real-time, minutes, hours?
-

## Problem Context & Why this is Hard

After clarifying, demonstrate you understand the problem's complexity:

This is a challenging system design problem for three reasons:

- **Scale:** 100K repositories generate roughly 10 million distinct functions. A brute-force search is too slow, we need sub-linear search complexity.
  - **Freshness:** Code changes constantly. The index must update within minutes of a git push without rebuilding the entire database.
  - **Precision:** Unlike chatbots, code search cannot hallucinate. If a user asks for Python code, returning JavaScript is a failure. We need high precision with language-specific filtering.
- 

## System Requirements (state these explicitly)

Requirement	Target	Rationale
Scale	100K repos → ~10M functions	Enterprise-scale codebase
Latency	P95 < 500ms	Developer productivity
Accuracy	Recall@5 > 80%, MRR > 0.6	Relevant results in top positions
Cost	~\$3K/month	Reasonable infrastructure spend
Availability	99.9% uptime	Business-critical tool


---

# Architecture Overview


At a high level, this system has two main pipelines:

- An **offline indexing pipeline** that processes code repositories and stores embeddings
- An **online query pipeline** that handles developer queries in real-time.

**Draw this diagram while explaining:**




Indexing Pipeline (Offline)



Query Pipeline (Online)

# Section 1: Indexing Pipeline (Offline)

Let me walk through the indexing pipeline first, since this is what populates our search index.



Indexing Pipeline (Offline)

## 1.1 Code Parsing

**What to say:** We use [Tree-sitter for AST](#) parsing because it is accurate, supports 40+ languages, and processes 1000 files/second.

**What we extract per function (example):**

```
{  
  "function_name": "authenticate_oauth2",  
  "signature": "def authenticate_oauth2(client_id: str, ...)",  
  "docstring": "Authenticate user using OAuth2 flow...",  
  "body": "import requests\n...",  
  "language": "python",  
  "file_path": "auth/oauth2.py",  
  "repo_name": "auth-library",  
  "stars": 1250,  
  "last_commit": "2024-01-15",  
  "imports": ["requests", "urllib.parse"]  
}
```

**Trade-off to mention:** We could use regex-based parsing which has no dependencies, but it's less accurate for complex code structures.

## 1.2 Chunking Strategy

**What to say:** We chunk at the function level because it provides semantic completeness, a whole function is easier to understand and use than arbitrary text chunks.

Approach	Chunk Size	Pros	Cons
Function-level (recommended)	200-500 tokens	Semantic completeness	Variable sizes
Fixed-size	512 tokens	Consistent	May split functions
File-level	Varies	Full context	Too large for embeddings

#### Special cases to mention:

- **Large functions (>500 tokens):** Split by logical blocks
- **Classes:** Index class + each method separately
- **Scripts:** Chunk by logical sections

**Numbers to know:** 10M functions × 350 average tokens = ~3.5B tokens to embed

---

## 1.3 Embedding Generation

**What to say:** For code embeddings, I recommend **CodeBERT** or **StarEncoder** because they are trained on code repositories and understand syntax and semantics. They are also free to self-host.

#### Embedding Model Comparison (At least memorize top 3-4):

Model	Dimensions	Best For	Cost
CodeBERT	768	General code	Free (self-host)
StarEncoder	768	Code search	Free (self-host)
OpenAI text-embedding-3-small	1536	High quality	~\$0.02/1M tokens
OpenAI text-embedding-3-large	3072	Highest quality	~\$0.65/1M tokens
Cohere embed-v3	1024	Code + text hybrid, multilingual	~\$0.10/1M tokens
Google text-embedding-004	768	General purpose, latest Gemini embedding	~\$0.025/1M tokens
Google text-multilingual-embedding-002	768	Multilingual support	~\$0.025/1M tokens
Voyage Code-2	1536	Code-specialized	~\$0.12/1M tokens
Jina AI v2	768	8K context, bilingual	~\$0.02/1M tokens
Nomic embed-text-v1.5	768	Long context (8K), open source	Free

### Batch processing numbers:

- **Batch size:** 32 functions
  - **GPU:** NVIDIA A100 (40GB)
  - **Throughput:** 10K functions/hour
  - **Total time for 10M functions:** ~42 days on 1 GPU → **4 days with 10 GPUs**
- 

## 1.4 Storage Architecture

**What to say:** We need three types of storage:

- A **Vector Database** for similarity search
- A **Relational Database** for metadata filtering
- And **Object Storage** for full source files.

Storage Type	Technology	Purpose	Monthly Cost
Vector DB	Pinecone/Qdrant	10M vectors, <50ms search	~\$500
Metadata DB	PostgreSQL	Filtering, boosting	~\$100
Code Storage	S3/GCS	Full source files	~\$10

**Vector DB selection (be ready to discuss trade-offs):**

Database	Latency	Best For	Self-Host?
Pinecone	<50ms	Zero-ops, enterprise SLAs	No
Qdrant	Lowest	Best price-performance	Yes
Weaviate	23-34ms	Complex data relationships	Yes
pgvector	10-100ms	Already using PostgreSQL	Yes
Milvus	Sub-10ms	Billion-scale workloads, need GPU acceleration, full OSS control	Yes

---

## 1.5 Incremental Indexing

**What to say:** Re-indexing 100K+ repos daily would be prohibitively expensive. Instead, we use incremental updates, only re-indexing files that changed since the last index.


```
# Pseudocode for incremental indexing
for repo in repos:
    if latest_commit > last_indexed_commit:
        changed_files = get_changed_files(repo, last_indexed_commit)
        for file in changed_files:
            delete_embeddings(repo, file)
            add_embeddings(repo, file)
```

**Key numbers:**

- **Daily changes:** ~1% of repos (1K repos)
  - **Re-indexing time:** 1 hour/day vs 4 days for full reindex
  - **Cost savings:** 96%
-

## Section 2: Query Pipeline (Online)

Now let me walk through what happens when a developer submits a query.



### 2.1 Query Understanding

**What to say:** Before we search, we need to understand the query intent and extract key entities.

**Input:** How do I authenticate with OAuth2 in Python?

**Output:**

```
{  
    "intent": "how_to",  
    "entities": ["OAuth2", "authentication"],  
    "expanded_terms": ["oauth", "auth", "token"],  
    "language_hint": "python",  
    "query_type": "code_example"  
}
```

**Techniques:**

1. **Intent Classification:** "how to" → implementation, "what is" → explanation
2. **Entity Extraction:** NER for libraries, frameworks, patterns
3. **Query Expansion:** Add synonyms ("auth" → "authentication")
4. **Language Detection:** Extract language filter from query

**Latency:** ~50ms (can use cached LLM call or lightweight classifier)

## 2.2 Multi-Stage Retrieval

**What to say:** This is the heart of the system. We use a three-stage retrieval process to balance speed and accuracy.

### Stage 1: Vector Similarity Search

- Embed query using same model as indexing
- Search vector DB for top 100 candidates
- Apply hard filters (language, minimum stars)
- **Latency:** ~50ms, Recall@100: ~95%

### Stage 2: Metadata Filtering & Boosting

- **Language match:** +20% score
- **Recency:** -10% per year old
- **Popularity (stars > 100):** +10%
- **Has docstring:** +15%
- **Latency:** ~20ms, Precision: 60% → 75%

### Stage 3: Cross-Encoder/flashrank Reranking

- Take top-20 from Stage 2
- Use cross-encoder or flashrank to score (query, code) pairs together
- Cross-encoders are more accurate than bi-encoders but slower
- **Latency:** ~100ms (20 calls × 5ms), MRR: 0.6 → 0.75

The key insight is that **Stage 1** optimizes for recall (don't miss relevant results), while **Stages 2 & 3** optimize for precision (rank the best results highest).

---

## 2.3 Response Generation

**What to say:** Once we have the top-5 code snippets, we assemble them into a context and use an LLM to generate a helpful response.

### Context Assembly:

Example 1 - authenticate\_oauth2 (python)  
Repository: auth-library (1250 ⭐)  
File: auth/oauth2.py  
[code snippet]

Example 2 - ...

### LLM Selection (know a few options):

Model	Input Cost	Output Cost	Best For
Gemini 2.5 Flash	\$0.30/1M	\$2.50/1M	Cost-efficient, 1M context
Claude Sonnet 4	\$3/1M	\$15/1M	Best for coding
GPT-4.1 Mini	\$0.40/1M	\$1.60/1M	Balanced

**Latency:** ~250ms for LLM generation

---

## 2.4 Caching Strategy

**What to say:** We use multi-level caching to reduce latency and cost for popular queries.

Cache Level	What's Cached	TTL	Hit Rate	Latency Saved
L1: Query Cache	Exact query → response	1 hour	~30%	495ms
L2: Embedding Cache	Query → embedding	24 hours	~50%	50ms
L3: Retrieval Cache	Embedding → top-100	6 hours	~40%	150ms

### Impact calculation:

- **Cache hit:** 5ms latency
  - **Cache miss:** 500ms latency
  - **Effective latency:**  $0.3 \times 5\text{ms} + 0.7 \times 500\text{ms} = 351\text{ms average}$
  - **Cost savings:** ~30% (fewer LLM calls)
-

## Section 3: Evaluation Metrics

Let me explain how we would measure success for this system.

### 3.1 Retrieval Metrics

MRR (Mean Reciprocal Rank): Position of first relevant result

- Target: **MRR > 0.6** (relevant result in top-2 on average)

Recall@K: Fraction of relevant docs found in top-K

- Target: **Recall@5 > 0.80**

Precision@K: Fraction of top-K that are relevant

- Target: **Precision@5 > 0.60**

### 3.2 Generation Metrics

**Human Evaluation** (sample 100 queries/week):

- **Helpfulness**: Does it solve the problem?
- **Correctness**: Is the code valid?
- **Clarity**: Easy to understand?

**User Feedback**: Target **thumbs-up rate > 75%**

### 3.3 Latency Metrics

Stage	Target (Milliseconds)
Query Understanding	50ms
Vector Search	50ms
Metadata Filtering	20ms
Reranking	100ms
LLM Generation	250ms
<b>Total P95</b>	<b>&lt; 500ms</b>

## Section 4: Scaling Strategy

### 4.1 Horizontal Scaling

Stateless services (easy to scale):


- **Query Service:** 10 instances
- **Embedding Service:** 5 instances
- **Reranking Service:** 5 instances

Autoscaling rules:

- Scale query service when **CPU > 70%**
- Scale embedding service when **queue > 100**
- Scale down nights/weekends

### 4.2 Database Sharding

**What to say:** We can shard the vector database by language, which gives us faster queries (smaller index per shard) and the ability to scale specific languages independently.



### 4.3 Cost Breakdown

Monthly costs (1M queries):

Component	Cost
Vector DB (Pinecone)	\$500
Metadata DB (PostgreSQL)	\$100
LLM (Gemini Flash)	\$1,645
Caching (Redis)	\$50
Compute (10 instances)	\$500
Storage (S3)	\$10
<b>TOTAL</b>	<b>\$2,805</b>

**Cost per query:** \$2,805 / 1M = **\$0.0028**

## Section 5: Debugging Poor Results

**Interview tip:** This is a common follow-up question. Have a systematic approach ready.

### Step 1: Measure Current Performance

```
results = evaluate_on_testset() # 100 queries with labels  
# Recall@5: 0.65, MRR: 0.52, Thumbs up: 58%
```

### Step 2: Categorize Failure Modes

Collect **thumbs down** examples and categorize:

Failure Mode	Percentage	Example
Wrong language	30%	Asked Python, got JavaScript
Outdated code	25%	Code from 2018
Too generic	20%	Not specific enough
Wrong intent	15%	Asked "what is", got "how to"
Bad explanation	10%	Good code, poor LLM response

**Key insight:** Pareto principle → fixing top 2-3 issues addresses 75% of problems.

### Step 3: Fix Biggest Issue

For wrong language (30% of failures):

- **Root cause:** Soft boosting wasn't strong enough
- **Solution:** Hard filtering by language when specified
- **A/B test result:** Recall@5 improved from 0.65 → 0.78

### Step 4: Monitor Improvement

Track metrics weekly and iterate.

## Section 6: Advanced Optimizations

Mention these if you have time or interviewer asks

### 6.1 Hybrid Search (Vector + BM25)

Combine semantic similarity with keyword matching:

```
final_score = 0.7 * vector_score + 0.3 * bm25_score
```

**Impact:** Recall@5: 0.80 → 0.85

### 6.2 Personalization

Boost results matching user's primary language and recent topics.

**Impact:** Precision: 0.60 → 0.70

### 6.3 Continuous Learning

Use clicks feedback to retrain reranker.

**Impact:** 2-3% quality improvement per month

---

## Interview Talking Points (At least Memorize these)

### When discussing scale

For 100K repos with 10M functions, I would use **Pinecone** for vector storage and **StarEncoder** for embeddings. The key is incremental indexing, only re-indexing changed files saves 96% of compute. With proper sharding by language and 3-tier caching, we can maintain [P95 latency](#) under 500ms while keeping costs at \$0.003 per query.

### When discussing accuracy

The secret to high accuracy is multi-stage retrieval. **Stage 1** uses vector search to get top-100 candidates with high recall. **Stage 2** applies metadata filters and boosting for language, recency, and popularity. **Stage 3** uses a cross-encoder to rerank the top-20, which improves MRR from 0.6 to 0.75. This three-stage approach balances speed and accuracy.

### When discussing debugging

When users report poor results, I follow a systematic process: measure current metrics on an eval set, categorize failure modes from user feedback, fix the biggest issue, A/B test the

fix, and monitor improvement weekly. This data-driven approach improved our thumbs-up rate from 58% to 75%.

## When discussing cost

At 1M queries/month, the main costs are LLM (\$1,645) and vector DB (\$500). We optimize by caching popular queries (30% hit rate = 30% savings), using Gemini Flash instead of GPT-4 (3x cheaper), and self-hosting the embedding model.

**Total cost:** \$0.003 per query.

---

## Quick Reference Numbers (For reference)

Metric	Value
Total Functions	10M
Vector Dimensions	768
Chunk Size	200-500 tokens
Vector Search Latency	50ms
Reranking Latency	100ms
LLM Latency	250ms
<b>Total P95 Latency</b>	<500ms
<b>Cost per Query</b>	<b>\$0.003</b>
Target MRR	>0.6
Target Recall@5	>0.80
Monthly Cost (1M queries)	~\$2,800

# Production Checklist

## Infrastructure

- Vector DB with replication
- Metadata DB with read replicas
- Redis cache cluster
- Load balancer
- Autoscaling
- S3 for code storage

## Monitoring

- Latency (P50/P95/P99)
- Error rate alerts (>1%)
- Cache hit rate
- LLM token usage
- Cost dashboard

## Quality

- Eval set with labels
- Weekly quality reports
- User feedback collection
- A/B testing framework

## Security

- API rate limiting
- Input validation
- Output sanitization
- Secrets management
- DDoS protection

---

Happy Learning 😊